

Computer Network Lab: Lab #1

塗大為

Department of Computer Science and Information Engineering
National Taiwan University
b07902024@ntu.edu.tw

陳威翰

Department of Computer Science and Information Engineering
National Taiwan University
b07902132@ntu.edu.tw

黃子軒

Department of Computer Science and Information Engineering
National Taiwan University
b07902134@ntu.edu.tw

楊子平

Department of Computer Science and Information Engineering
National Taiwan University
b07902136@ntu.edu.tw

I. INSTALLATION & EXECUTION

Run `make` in the top-level directory to compile the `traceroute` executable. The program supports a subset of options implemented by Linux `traceroute`¹:

- `-f` Set the TTL of the first packet. (Default: 1)
- `-m` Set the maximum TTL in the sent packets. (Default: 30)
- `-w` The waiting time (in seconds) when receiving packets. (Default: 5)
- `-q` The number of queries for a single hop. (Default: 3)

Note that, to mimic the behavior of Linux `traceroute`, the default protocol used is UDP instead of ICMP. This can be toggled with the following options:

- `-I` ICMP mode
- `-T` TCP mode

II. IMPLEMENTATION

In general, for each hop, the `traceroute` program sends several packets (depending on the option `-q`) with a fixed TTL set in the IP header. The packets either reach the destination or get discarded half-way through due to the insufficient TTL. The former case should be handled differently for each protocol, which we will cover in the following paragraphs. In the latter case, because the routers send ICMP replies when discarding packets, regardless of the transmission layer protocol, the same handling strategy can be shared among different modes.

Specifically, our program waits for two types of ICMP reply after sending requests in all protocols: destination

unreachable (type 3) and time exceeded (type 7). When encountering the former reply, we simply print out `!h`, `!n` or `!p` depending on the corresponding ICMP code. The latter reply suggests that the sent packet does not have enough TTL to proceed to the destination, and thus we increment the hop counter and continue tracing. Additionally, since the routers might accidentally discard our packets without notifying us, the socket used for receiving reply is configured with a certain timeout (specified by the `-w` option). When `recvfrom()` times out, `*` will be printed.

For ICMP, we simply issue ICMP echo requests to the destination. When the destination is reached, the server will reply to us an ICMP echo reply. The reply is then verified by comparing the identifier and sequence number with the expected values.

For UDP, we incorporate the same strategy as Linux `traceroute`. In particular, our program sends a UDP datagram with a fixed TTL to a specific port, which starts from an unusual number of 33435 and increments after each probe. Because the port we target is rarely used, we can assume that the destination is reached upon receiving ICMP port-unreachable messages and terminate the program. Additionally, the reply can be easily verified by the destination port of the UDP datagram it wraps.

Finally, for TCP, we use `connect` directly —again with a fixed TTL—to connect to the destination on port 80. Non-blocking mode together with `select` are used to ensure the timeout is honored. We then assume a packet has reached the destination if the connection is accepted or an ICMP port-unreachable message is received. Similar to UDP, replies are verified by their destination ports.

¹Throughout, this refers to Dmitry Butskoy's implementation.

III. QUESTIONS

A. Detecting and Defending against *traceroute*

A possible way of detecting *traceroute* is to have routers that the server operator control in front of the server itself. Packets with low TTL can then be logged and correlated (e.g., if the same source sends multiple such packets) to spot potential *traceroute* attempts.

As for defending against *traceroute*, it is difficult to see in what threat model this actually makes sense. One can already obtain a lot of information from the destination IP address (e.g., AS number²). If one desires to obfuscate such details, it is likely that she also wishes to hide the destination IP by the use of CDNs or even Tor hidden services.

However, if one still wishes to hide the *traceroute* details, one method is to have routers that do not respond with ICMP time-exceed messages in front of the server.

B. The Reason *traceroute* Cannot Show the Full Route

Some intermediate routers may choose not to reply with ICMP time-exceed messages. Alternatively, sometimes a router is busy and may not have the resources to send out such messages. In these cases, the path cannot be shown by the *traceroute* command.

C. Why the Result May Not Always Be the Same

The route a packet takes may depend on the underlying algorithms used by the intermediate routers. For instance, to perform load-balancing, a router might forward packets to different destinations even if they are identical. This may result in inconsistent outputs from the *traceroute* program.

IV. COMPARISON

A. Domestic and Foreign

All the following experiments are conducted in the CSIE building, using the WiFi access points csie{-5G} and the *traceroute* program we implemented running in ICMP mode.

The major difference between tracerouting a domestic server and a foreign server is the length of the route and the overall response time. Quite obviously, tracerouting foreign servers generally take more hops and more time than domestic servers.

Furthermore, packets that are bounded for foreign servers are likely to pass through *gateways* that coordinate packets leaving Taiwan. For example, when tracerouting codeforces.com, the round-trip time increases from 50ms to 200ms after hopping over the routers at 203.78.181.38, as can be seen in Figure 1. Using ipinfo.io, we can confirm that this IP address belongs to the autonomous system Taiwan Internet Gateway, a subsidiary of Chunghua Telecom that specializes in international peering.

²Which can lead to additional peering information.

```

9 173.61.41.175 (4.334 ms)
10 203.78.181.38 (57.786 ms)
11 sngc-bl-link.ip.twelvet99.net (205.289 ms)
12 sngc-b3-link.ip.twelvet99.net (196.343 ms)

```

Figure 1. Taiwan Internet Gateway

The detailed results can be found in Section V-A.

B. TCP, UDP, and ICMP

An experiment was conducted to traceroute codeforces.com (213.248.110.126). The results can be found in Section V-B.

From this and several other experiments, we find that some destinations silently drop UDP but not TCP and ICMP packets, making the results of the latter two more useful in these cases.

We also observed that TCP and UDP were more sensitive to the selection of destination ports. For example, if we send UDP datagrams to a more commonly used UDP port, say 53, the requests are more likely to be dropped. In TCP, choosing ports other than 80 can similarly mess up the results.

This phenomenon can be seen from an example performing a *traceroute* to google.com, which is shown in Section V-C.

V. APPENDIX

A. Experiment 1

- Domestic
 - CSIE website (140.112.30.26)
 - 1 10.5.7.253 (10.5.7.253) 6.936 ms
 - 3.947 ms 8.016 ms
 - 2 www.inm.ntu.edu.tw (140.112.30.26)
 - 1.173 ms 1.110 ms 1.119 ms
 - TSMC (202.126.64.15)
 - 1 LPC.mshome.net (172.18.91.161) 0.779 ms
 - 0.244 ms 0.282 ms
 - 2 10.5.7.253 (10.5.7.253) 9.459 ms
 - 89.653 ms 56.104 ms
 - 3 172.17.0.2 (172.17.0.2) 29.228 ms
 - 3.880 ms 8.138 ms
 - 4 140.112.16.190 (140.112.16.190) 12.003 ms
 - 15.971 ms 13.277 ms
 - 5 140.112.149.121 (140.112.149.121)
 - 10.287 ms 32.065 ms 15.863 ms
 - 6 140.112.0.242 (140.112.0.242) 70.729 ms
 - 5.701 ms 7.289 ms
 - 7 140.112.0.206 (140.112.0.206) 11.952 ms
 - 21.154 ms 8.999 ms
 - 8 tpd-3301.hinet.net (211.22.226.202)
 - 14.342 ms 21.845 ms 9.155 ms
 - 9 TPD-3012.hinet.net (220.128.4.150)
 - 6.501 ms 44.936 ms 15.528 ms
 - 10 220-128-25-138.HINET-IP.hinet.net (220.128.25.138) 16.592 ms
 - 28.574 ms 31.653 ms

11 sczs-3312.hinet.net (220.128.2.57)
 13.017 ms 20.755 ms 24.282 ms
 12 h145.s222.ts.hinet.net (168.95.222.145)
 8.897 ms 26.827 ms 82.649 ms
 13 202-39-252-25.HINET-IP.hinet.net
 (202.39.252.25) 9.079 ms 86.673 ms
 22.757 ms

- Foreign

- codeforces.com (87.240.191.209)
 - 1 LPC.mshome.net (172.18.91.161) 1.236
ms 0.497 ms 0.381 ms
 - 2 10.5.7.253 (10.5.7.253) 28.652 ms
12.865 ms 36.669 ms
 - 3 172.17.0.2 (172.17.0.2) 8.931 ms
13.525 ms 1.961 ms
 - 4 140.112.16.190 (140.112.16.190) 17.624
ms 8.226 ms 11.834 ms
 - 5 140.112.149.121 (140.112.149.121)
8.882 ms 51.239 ms 13.968 ms
 - 6 140.112.0.222 (140.112.0.222) 131.883
ms 5.461 ms 3.388 ms
 - 7 140.112.0.206 (140.112.0.206) 7.022 ms
9.755 ms 13.383 ms
 - 8 203.160.226.233 (203.160.226.233)
32.465 ms 9.696 ms 31.464 ms
 - 9 173-61-41-175.TWGATE-IP.twgate.net
(175.41.61.173) 15.454 ms 20.913 ms
12.009 ms
 - 10 203.78.181.38 (203.78.181.38) 60.726
ms 53.249 ms 50.882 ms
 - 11 snge-b1-link.ip.twelve99.net
(62.115.33.244) 195.887 ms 264.716 ms
202.052 ms
 - 12 snge-b3-link.ip.twelve99.net
(62.115.124.86) 228.950 ms 194.462 ms
197.459 ms
 - 13 mei-b5-link.ip.twelve99.net
(62.115.124.123) 254.199 ms 273.785
ms 262.222 ms
 - 14 ffm-bb2-link.ip.twelve99.net
(62.115.124.60) 270.268 ms 275.339 ms
292.786 ms
 - 15 s-bb2-link.ip.twelve99.net
(62.115.138.104) 272.092 ms 292.729
ms 287.145 ms
 - 16 sap-b4-link.ip.twelve99.net
(62.115.115.247) 289.984 ms 336.039
ms 278.236 ms
 - 17 m100-svc072762-ic360359.ip.twelve99-
cust.net (62.115.145.97) 277.983 ms
279.961 ms 267.874 ms
 - 18 srv209-191-240-87.vk.com
(87.240.191.209) 291.344 ms 294.120
ms 289.443 ms

- poj.org (168.235.95.213)

- 1 _gateway (10.5.7.253) 156.067 ms
87.602 ms 10.674 ms
- 2 172.17.0.2 (172.17.0.2) 2.559 ms
1.984 ms 12.833 ms
- 3 140.112.16.190 (140.112.16.190) 10.694
ms 11.297 ms 7.015 ms
- 4 140.112.149.121 (140.112.149.121)
3.051 ms 3.769 ms 3.706 ms
- 5 140.112.0.222 (140.112.0.222) 2.900 ms
4.139 ms 9.479 ms
- 6 140.112.0.206 (140.112.0.206) 9.862 ms
18.061 ms 2.800 ms

- 7 203.160.226.133 (203.160.226.133)
40.297 ms
203.160.226.233 (203.160.226.233)
4.374 ms
203.160.226.133 (203.160.226.133)
3.971 ms
- 8 181-61-41-175.TWGATE-IP.twgate.net
(175.41.61.181) 10.232 ms 3.271 ms
3.014 ms
- 9 230-60-41-175.TWGATE-IP.twgate.net
(175.41.60.230) 138.214 ms 173.492
ms 134.874 ms
- 10 38.122.183.73 (38.122.183.73) 148.181
ms 165.018 ms 154.132 ms
- 11 be2431.ccr41.sjc03.atlas.cogentco.com
(154.54.88.189) 159.868 ms 178.960 ms
157.573 ms
- 12 te0-3-0-4.128.br04.sjo01.pccwbbtn.net
(63.217.21.97) 157.422 ms 146.006 ms
141.384 ms
- 13 TenGE0-2-0-0.br03.lax04.pccwbbtn.net
(63.218.50.177) 145.887 ms 146.587 ms
152.753 ms
- 14 * * *
- 15 168.235.95.213 (168.235.95.213)
139.287 ms 138.260 ms 154.153 ms

B. Experiment 2

- UDP

- 1 _gateway (10.5.7.253) 5.177 ms 12.853
ms 7.370 ms
- 2 172.17.0.2 (172.17.0.2) 1.063 ms
9.192 ms 8.347 ms
- 3 140.112.16.190 (140.112.16.190) 14.703
ms 9.227 ms 4.747 ms
- 4 140.112.149.121 (140.112.149.121)
8.071 ms 8.947 ms 2.213 ms
- 5 140.112.0.222 (140.112.0.222) 4.105 ms
1.352 ms 15.247 ms
- 6 140.112.0.206 (140.112.0.206) 10.483
ms 2.529 ms 2.092 ms
- 7 203.160.226.133 (203.160.226.133)
17.977 ms 2.576 ms 9.485 ms
- 8 181-61-41-175.TWGATE-IP.twgate.net
(175.41.61.181) 8.921 ms
173-61-41-175.TWGATE-IP.twgate.net
(175.41.61.173) 2.519 ms 6.870 ms
- 9 203.78.181.174 (203.78.181.174) 56.258
ms
203.78.181.38 (203.78.181.38) 50.063
ms 51.124 ms
- 10 snge-b1-link.ip.twelve99.net
(62.115.33.244) 191.701 ms 286.606 ms
211.198 ms
- 11 snge-b3-link.ip.twelve99.net
(62.115.124.86) 192.443 ms 203.274 ms
212.645 ms
- 12 mei-b5-link.ip.twelve99.net
(62.115.124.123) 253.050 ms 249.095
ms
mei-b4-link.ip.twelve99.net
(62.115.143.22) 244.832 ms
- 13 ffm-bb2-link.ip.twelve99.net
(62.115.124.60) 268.867 ms
ffm-bb1-link.ip.twelve99.net
(62.115.116.20) 273.898 ms 273.515
ms

```

14 s-bb2-link.ip.twelve99.net
   (62.115.138.104) 278.886 ms
   s-bb3-link.ip.twelve99.net
     (62.115.138.236) 274.170 ms *
15 sap-b4-link.ip.twelve99.net
   (62.115.115.249) 291.637 ms 272.464
   ms 275.262 ms
16 m100-svc072762-ic360359.ip.twelve99-
   cust.net (62.115.145.97) 282.819 ms
   292.106 ms 275.618 ms
17 srv209-191-240-87.vk.com
   (87.240.191.209) 289.786 ms 279.066
   ms
   srv247-191-240-87.vk.com
     (87.240.191.247) 276.316 ms
18 * * *
19 * * *
20 * * *
21 * * *
22 * * *
23 * * *
24 * * *
25 * * *
26 * * *
27 * * *
28 * * *
29 * * *
30 * * *

```

• TCP

```

1 _gateway (10.5.7.253) 4.778 ms 8.129
   ms 9.314 ms
2 172.17.0.2 (172.17.0.2) 1.375 ms
   3.463 ms 7.266 ms
3 140.112.16.190 (140.112.16.190) 8.777
   ms 5.986 ms 4.053 ms
4 140.112.149.121 (140.112.149.121)
   2.971 ms 5.007 ms 2.182 ms
5 140.112.0.222 (140.112.0.222) 2.505 ms
   1.887 ms 3.493 ms
6 140.112.0.206 (140.112.0.206) 3.597 ms
   3.256 ms 2.335 ms
7 203.160.226.133 (203.160.226.133)
   3.368 ms
   203.160.226.233 (203.160.226.233)
     7.892 ms 2.289 ms
8 181-61-41-175.TWGATE-IP.twgate.net
   (175.41.61.181) 2.468 ms 3.031 ms
   173-61-41-175.TWGATE-IP.twgate.net
     (175.41.61.173) 2.524 ms
9 203.78.181.38 (203.78.181.38) 57.954
   ms 53.096 ms 60.469 ms
10 snge-b1-link.ip.twelve99.net
   (62.115.33.244) 194.190 ms 194.756 ms
   200.354 ms
11 snge-b3-link.ip.twelve99.net
   (62.115.124.86) 191.790 ms 193.862 ms
   203.713 ms
12 mei-b4-link.ip.twelve99.net
   (62.115.143.22) 243.668 ms 243.337 ms
   mei-b5-link.ip.twelve99.net
     (62.115.124.123) 248.419 ms
13 ffm-bb2-link.ip.twelve99.net
   (62.115.114.202) 265.831 ms
   ffm-bb1-link.ip.twelve99.net
     (62.115.116.20) 266.213 ms
   ffm-bb2-link.ip.twelve99.net
     (62.115.124.60) 269.276 ms

```

```

14 s-bb3-link.ip.twelve99.net
   (62.115.138.236) 274.976 ms
   s-bb2-link.ip.twelve99.net
     (62.115.138.104) 270.357 ms
   s-bb3-link.ip.twelve99.net
     (62.115.138.236) 273.353 ms
15 sap-b4-link.ip.twelve99.net
   (62.115.115.247) 279.866 ms 265.030
   ms 273.439 ms
16 m100-svc072762-ic360359.ip.twelve99-
   cust.net (62.115.145.97) 279.134 ms
   272.046 ms 270.782 ms
17 srv247-191-240-87.vk.com
   (87.240.191.247) 268.101 ms 287.006
   ms
   srv209-191-240-87.vk.com
     (87.240.191.209) 307.428 ms
18 * * *
19 213-248-110-126.teliacarrier-cust.com
   (213.248.110.126) 272.691 ms 289.466
   ms 274.597 ms

```

• ICMP

```

1 _gateway (10.5.7.253) 13.040 ms
   12.230 ms 6.248 ms
2 172.17.0.2 (172.17.0.2) 3.782 ms
   5.563 ms 10.574 ms
3 140.112.16.190 (140.112.16.190) 11.496
   ms 10.306 ms 8.886 ms
4 140.112.149.121 (140.112.149.121)
   7.881 ms 1.962 ms 6.497 ms
5 140.112.0.222 (140.112.0.222) 2.087 ms
   8.428 ms 1.921 ms
6 140.112.0.206 (140.112.0.206) 2.962 ms
   7.081 ms 2.648 ms
7 203.160.226.133 (203.160.226.133)
   7.442 ms 7.132 ms 13.198 ms
8 173-61-41-175.TWGATE-IP.twgate.net
   (175.41.61.173) 10.789 ms 3.117 ms
   2.813 ms
9 203.78.181.38 (203.78.181.38) 62.771
   ms 56.156 ms 54.199 ms
10 snge-b1-link.ip.twelve99.net
   (62.115.33.244) 203.994 ms 208.414 ms
   193.165 ms
11 snge-b3-link.ip.twelve99.net
   (62.115.124.86) 194.337 ms 222.604 ms
   195.796 ms
12 mei-b5-link.ip.twelve99.net
   (62.115.124.123) 250.144 ms 253.142
   ms 248.075 ms
13 ffm-bb1-link.ip.twelve99.net
   (62.115.124.58) 279.157 ms 280.623 ms
   283.093 ms
14 s-bb3-link.ip.twelve99.net
   (62.115.138.236) 293.509 ms 291.686
   ms 292.873 ms
15 sap-b4-link.ip.twelve99.net
   (62.115.115.249) 283.721 ms 289.526
   ms 283.743 ms
16 m100-svc072762-ic360359.ip.twelve99-
   cust.net (62.115.145.97) 281.130 ms
   284.971 ms 285.200 ms
17 srv247-191-240-87.vk.com
   (87.240.191.247) 285.636 ms 304.410
   ms 287.190 ms
18 * * *
19 213-248-110-126.teliacarrier-cust.com

```

```
(213.248.110.126) 289.667 ms 291.852
ms 290.824 ms
```

C. Experiment 3

- Destination port fixed to 53:

```
1 _gateway (192.168.40.1) 0.730 ms
  0.617 ms 0.599 ms
2 gateway243.m3.ntu.edu.tw
  (140.112.243.254) 1.312 ms 1.273 ms
  1.239 ms
3 140.112.0.98 (140.112.0.98) 1.074 ms
  1.185 ms 1.174 ms
4 140.112.0.250 (140.112.0.250) 1.009 ms
  1.180 ms 1.256 ms
5 140.112.0.206 (140.112.0.206) 3.106 ms
  1.530 ms 1.628 ms
6 140.112.0.34 (140.112.0.34) 1.919 ms
  1.733 ms 1.428 ms
7 72.14.196.229 (72.14.196.229) 1.909 ms
  2.162 ms 2.238 ms
8 * * *
9 * * *
10 * * *
[snip]
```

- Destination port counting up from 33435:

```
1 _gateway (192.168.40.1) 0.909 ms
  0.745 ms 0.704 ms
2 gateway243.m3.ntu.edu.tw
  (140.112.243.254) 1.519 ms 1.285 ms
  1.361 ms
3 140.112.0.98 (140.112.0.98) 1.380 ms
  1.321 ms 1.173 ms
4 140.112.0.230 (140.112.0.230) 1.300 ms
  140.112.0.250 (140.112.0.250) 1.088 ms
  1.314 ms
5 140.112.0.206 (140.112.0.206) 8.233 ms
  2.036 ms 1.774 ms
6 140.112.0.34 (140.112.0.34) 2.888 ms
  1.670 ms 1.835 ms
7 72.14.196.229 (72.14.196.229) 6.426 ms
  2.087 ms 2.333 ms
8 108.170.244.97 (108.170.244.97) 2.597
  ms
  108.170.244.129 (108.170.244.129)
  3.805 ms
  108.170.244.97 (108.170.244.97) 2.515
  ms
9 216.239.48.135 (216.239.48.135) 3.181
  ms
  108.170.225.177 (108.170.225.177)
  3.331 ms 3.346 ms
10 tsa03s06-in-f14.1e100.net
  (172.217.160.110) 2.277 ms 2.290 ms
  2.322 ms
```